


Baidu-III Project-II

AI for CV Group
2020



Week 7-8 II Face Detection & Recognition

Contents:

I. Face Detection

- A. Introduction
- B. A real example: FaceBoxes
- C. Code of FaceBoxes
- D. Advices for Projects

II. Face Recognition

- E. What Is Recognition
- F. Face Recognition Procedure
- G. A real example: LightCNN

I. Face Detection

I. Face Detection

A. Introduction:

Red: Have to know before interview

1998~

Development: **ASM** — Viola Jones — MTCNN — SSH — BlazeFace



- AAM / CLM
- PCA

I. Face Detection



A. Introduction:

Pink: Better to know before interview

Blue: Code!!

2001

Development: ASM — Viola Jones — MTCNN — SSH — BlazeFace



- Haar Feature
- Integral Image
- Adaboost
- Cascade

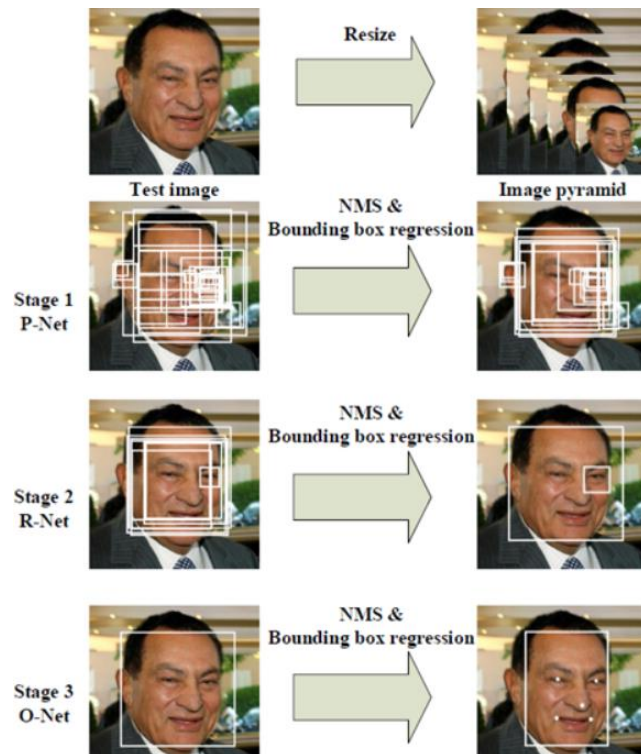
I. Face Detection

A. Introduction:

Pink: Better to know before interview

2015

Development: ASM — Viola Jones — MTCNN — SSH — BlazeFace



- Accurate and fast

- 3 stages / 3 networks

P-Net: Proposal Net

R-Net: Refined Net


O-Net: Output Net

I. Face Detection

A. Introduction:

Green: Good to know before interview

Development: ASM — Viola Jones — MTCNN — ²⁰¹⁷SSH — BlazeFace



- Accurate and fast
- Small faces in wild
- SSD-like structure

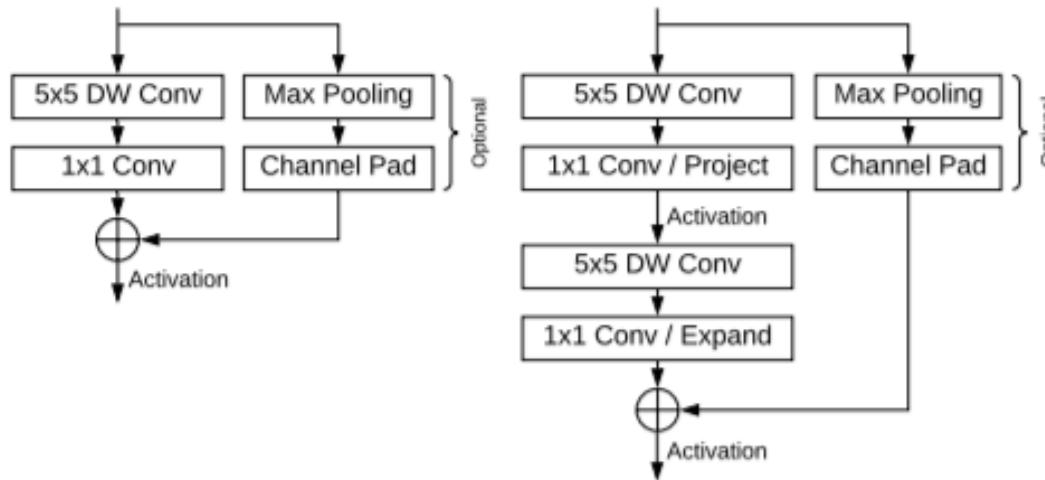
I. Face Detection

A. Introduction:

Green: Good to know before interview

2019

Development: ASM — Viola Jones — MTCNN — SSH — BlazeFace



- Accurate and extremely fast!
- Small faces in wild
- Latest tech in face detection
- Realtime in mobile devices

I. Face Detection

A. Introduction:

Trend: Small face in the wild
Real time in mobile devices
[Anti-spoof]

Companies: Megvii (Face++)
Sensetime
Google
Economic Companies
[Alibaba / Amazon / VISA / ...]

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

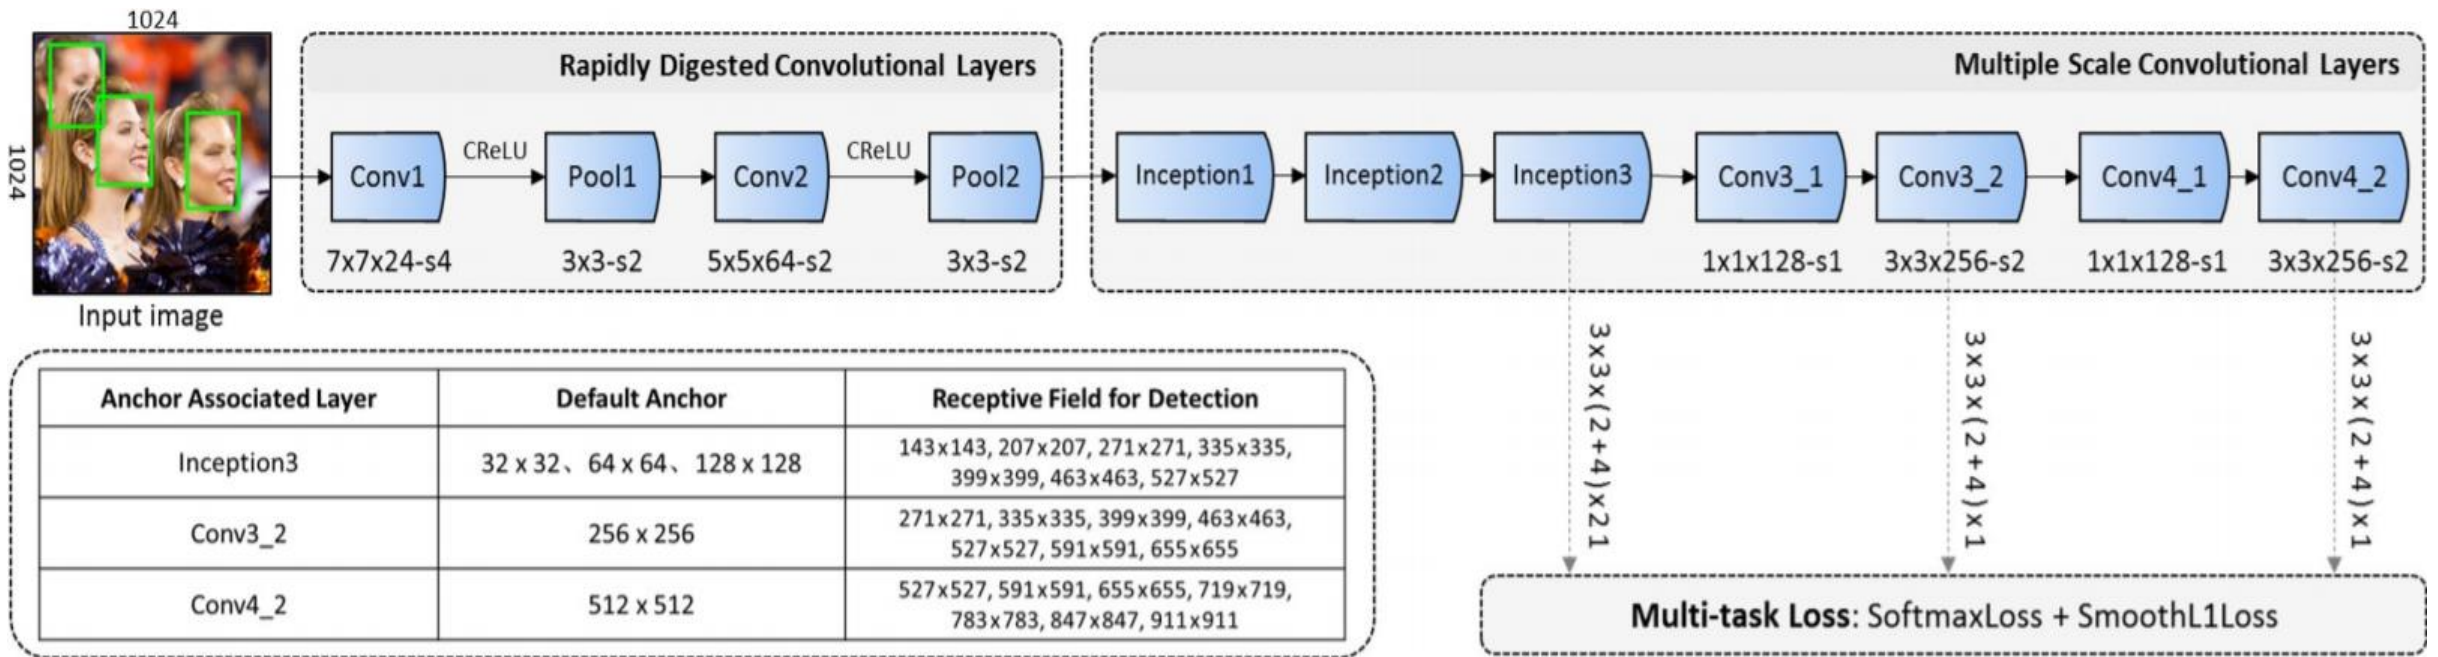
Features: Fast! Real-time on CPU
100+FPS on GPU
Practical & Acceptable



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

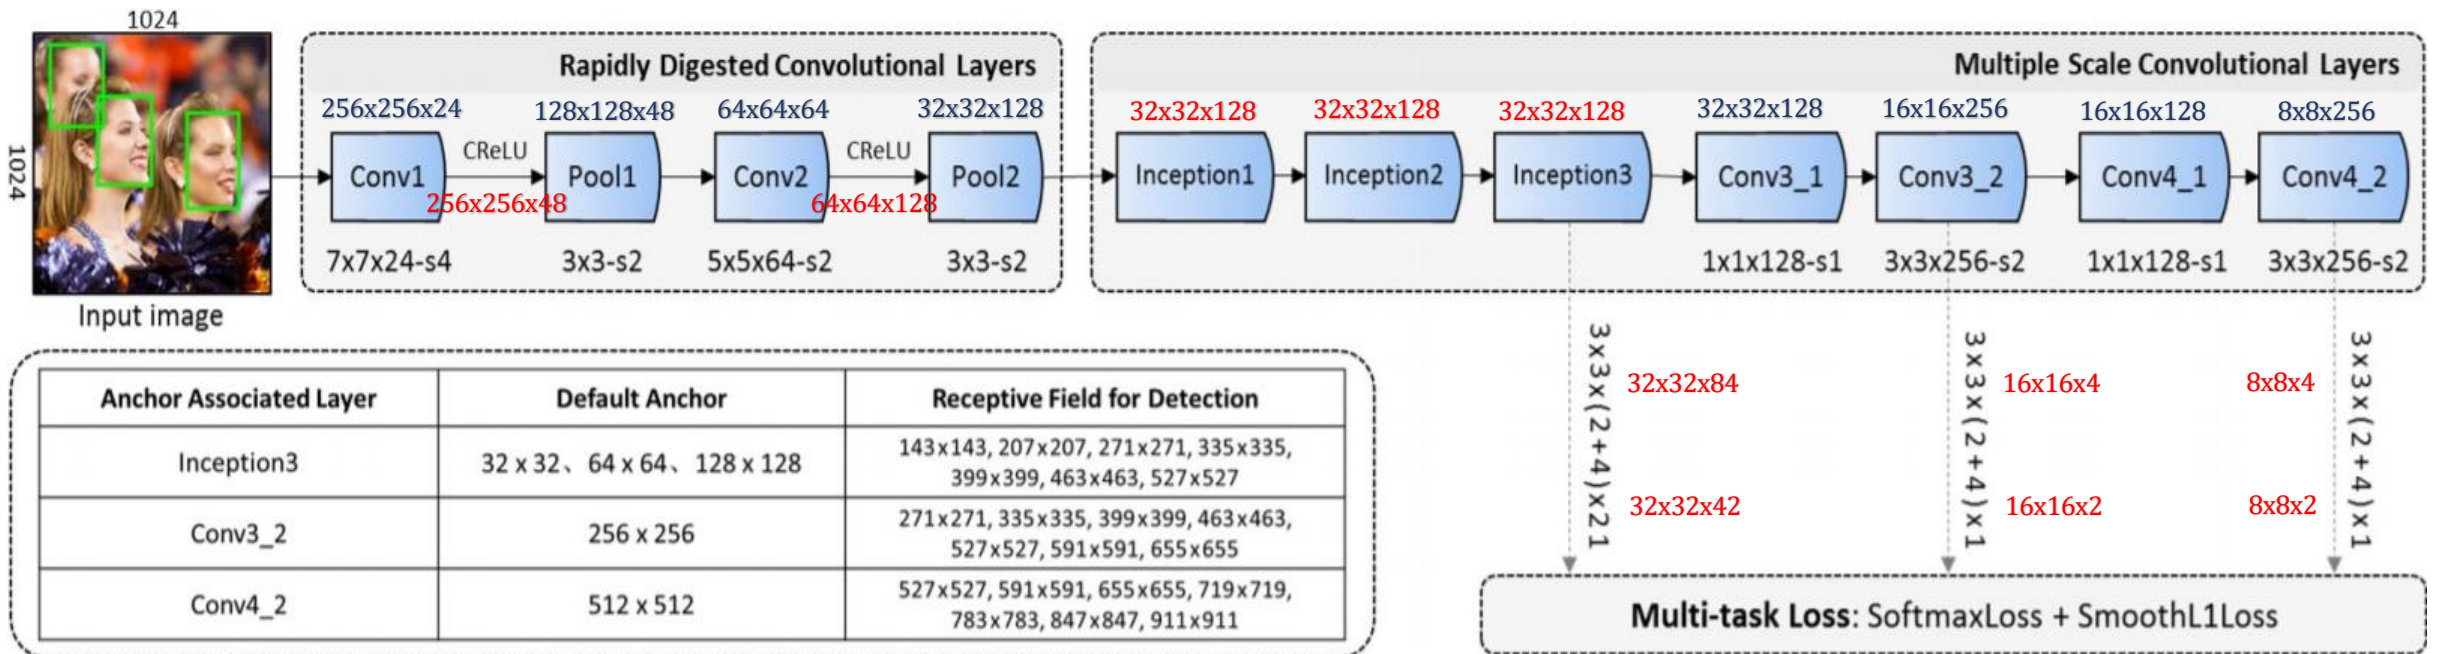
Structure:



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

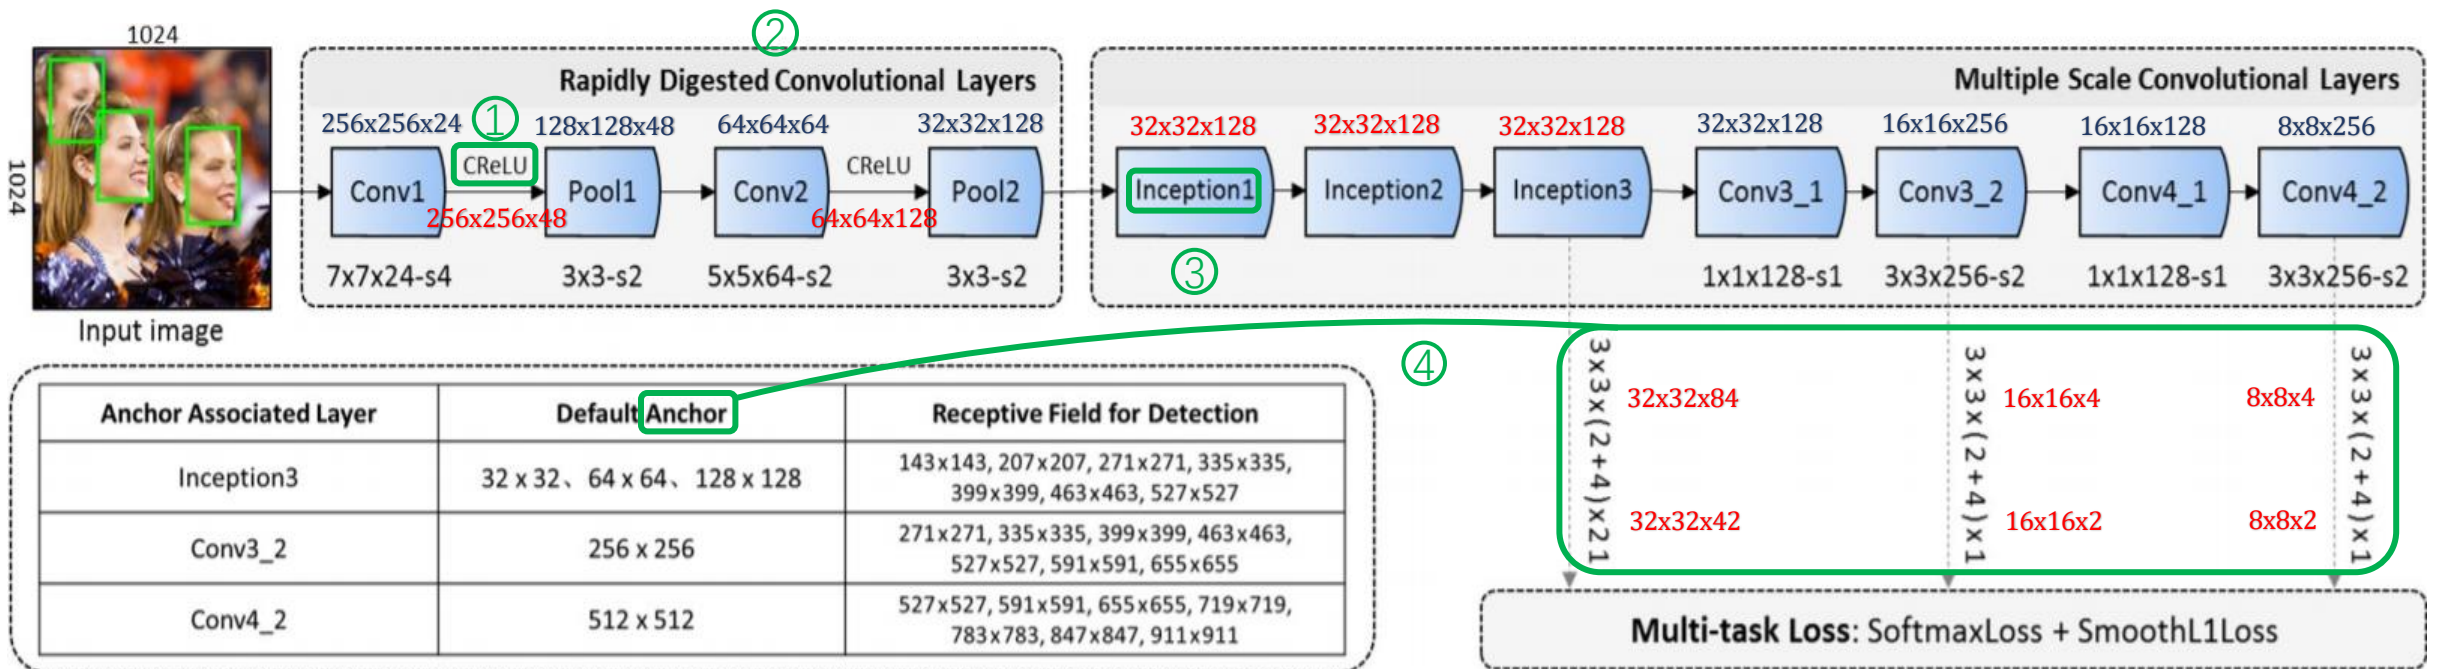
Structure:



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure:




I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 1. CRelu [2016]

```
class CReLU(nn.Module):  
    def __init__(self):  
        super(CReLU, self).__init__()  
    def forward(self, x):  
        return torch.cat((F.relu(x), F.relu(-x)), 1)
```

- Reduce parameters / Accelerate system: $\frac{n}{2}$ filters, n channels
- Filters of lower conv layers form pairs



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 1. CReLU [2016]


```
class CReLU(nn.Module):
    def __init__(self):
        super(CReLU, self).__init__()
    def forward(self, x):
        return torch.cat((F.relu(x), F.relu(-x)), 1)
```

- Don't be bewildered by CELU

CELU: continuously differentiable exponential linear units

CReLU: concatenated rectified linear units


- Reduce parameters / Accelerate system: $\frac{n}{2}$ filters, n channels
- Filters of lower conv layers form pairs



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 2. Rapid Digested Convolutional Layers


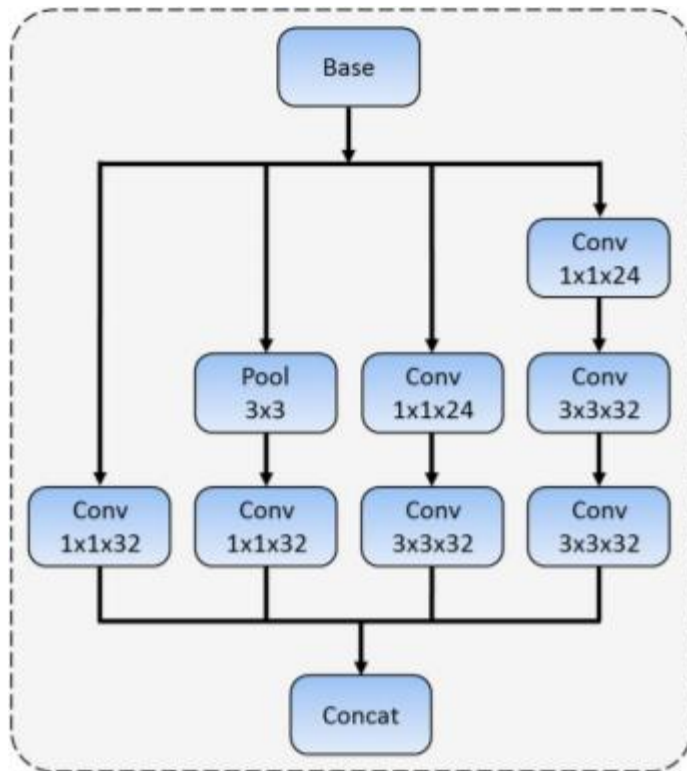


- Rapidly shrink the input size: 7x7
- Carefully choose filter size: bigger for getting info; smaller for speeding up
- Reducing channel number: CReLU

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 3. Inception




- **1x1 Conv:**
reduce dimension
- **Inception Module:**
multiple resolution

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:


Structure: 4. Anchor



I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 4. Anchor

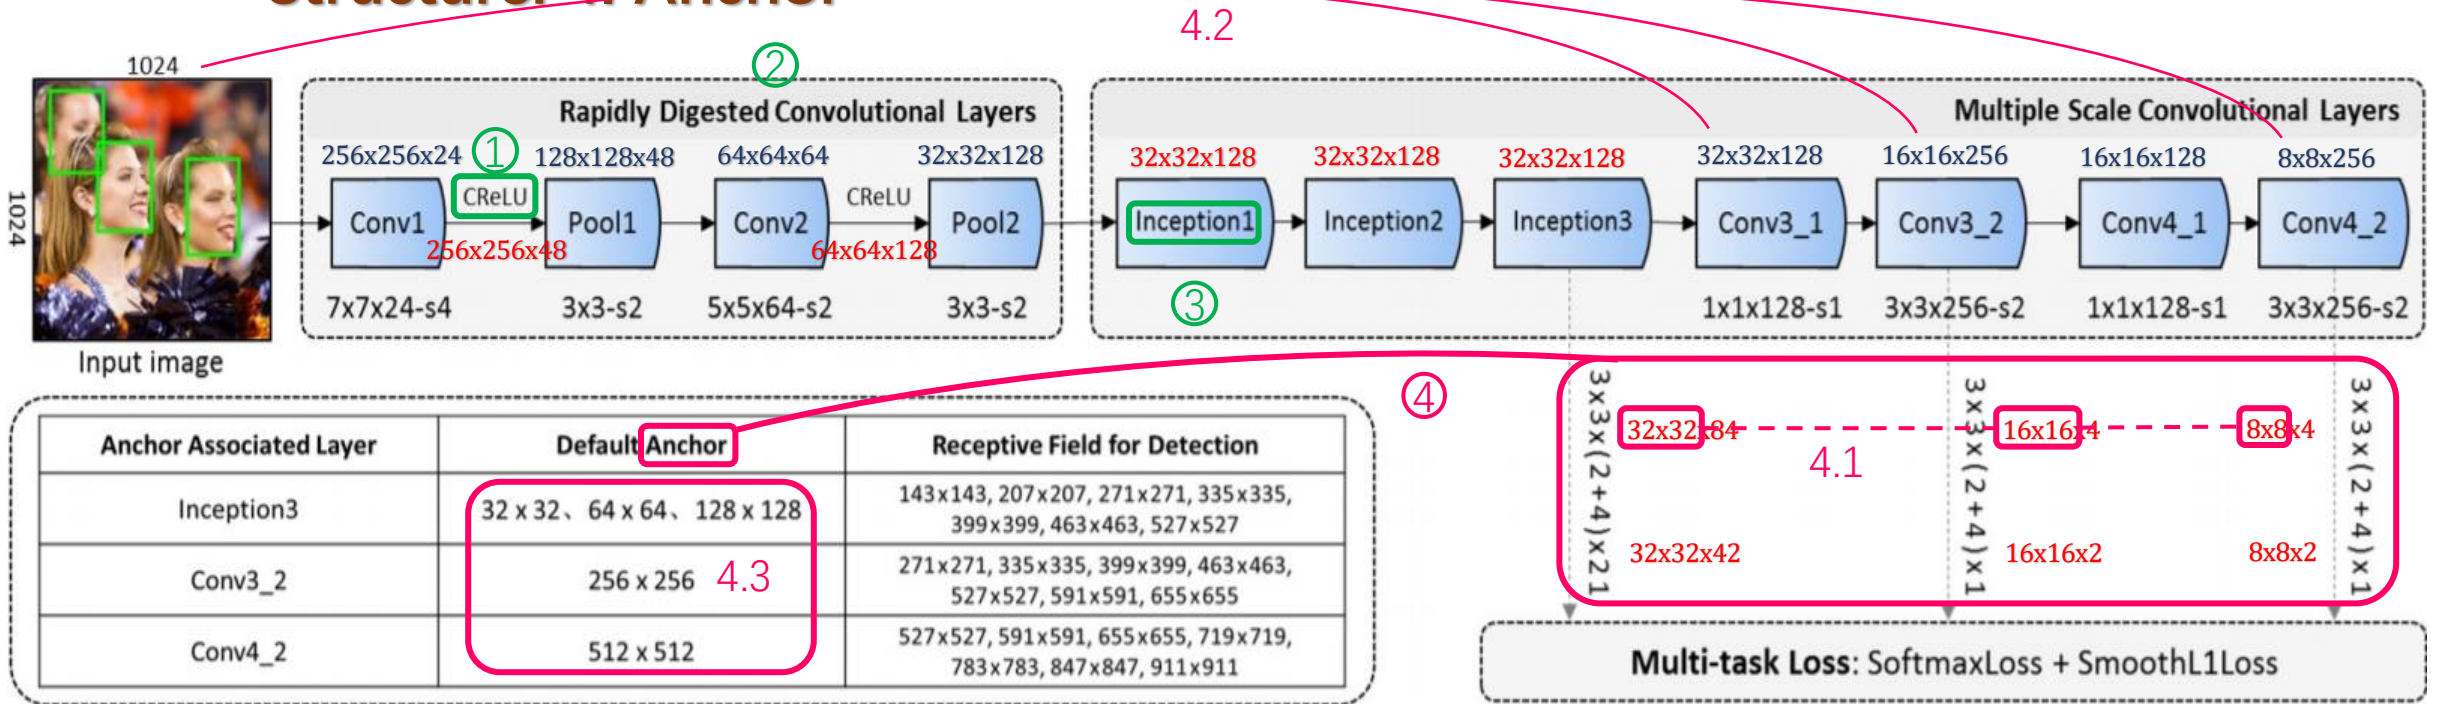


- 4.1 feature maps: 3 scales: 32 x 32, 16 x 16, 8 x 8

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 4. Anchor

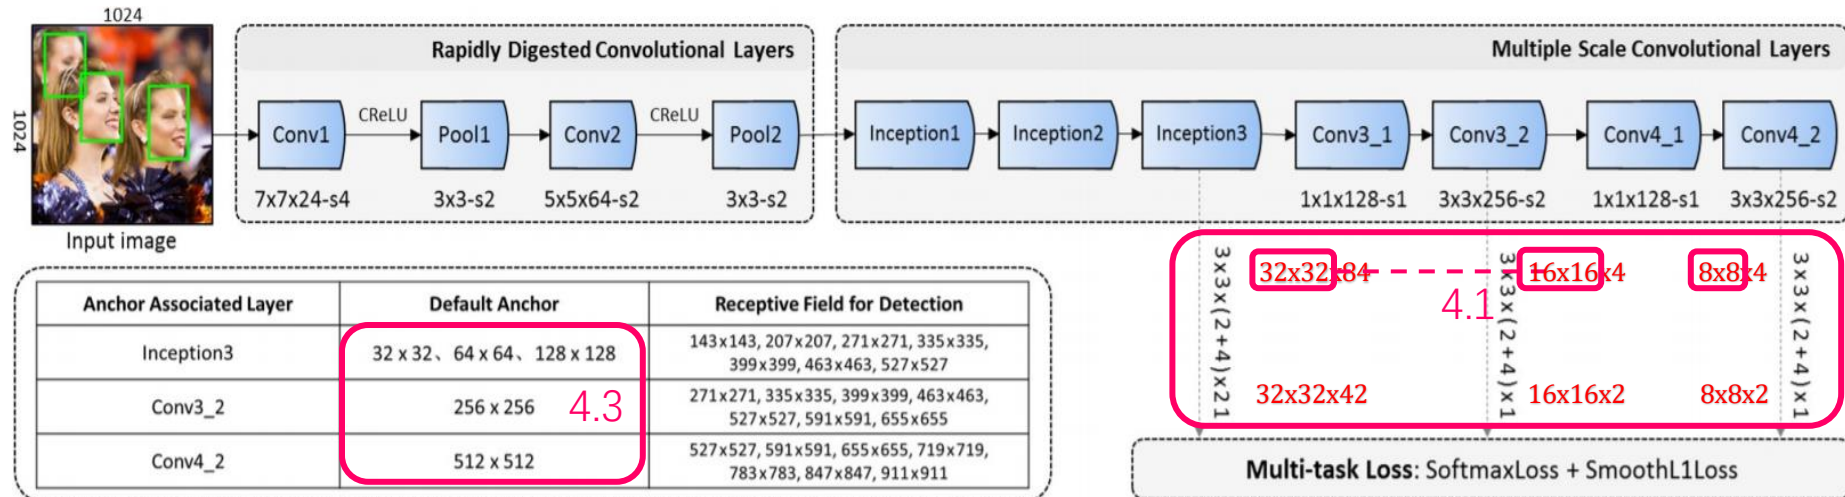


- 4.2 skips: $1024 / 32 = 32$; $1024 / 64 = 16$; $1024 / 128 = 8$

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 4. Anchor



- 4.3 min_sizes [anchor size]:

min_sizes	skips	feature maps
[32x32, 64x64, 128x128]	32	32x32
[256x256]	64	16x16
[512x512]	128	8x8

I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 4. Anchor

- 4.3 min_sizes
[anchor size]:

min_sizes	skips	feature maps
[32x32, 64x64, 128x128]	32	32x32
[256x256]	64	16x16
[512x512]	128	8x8

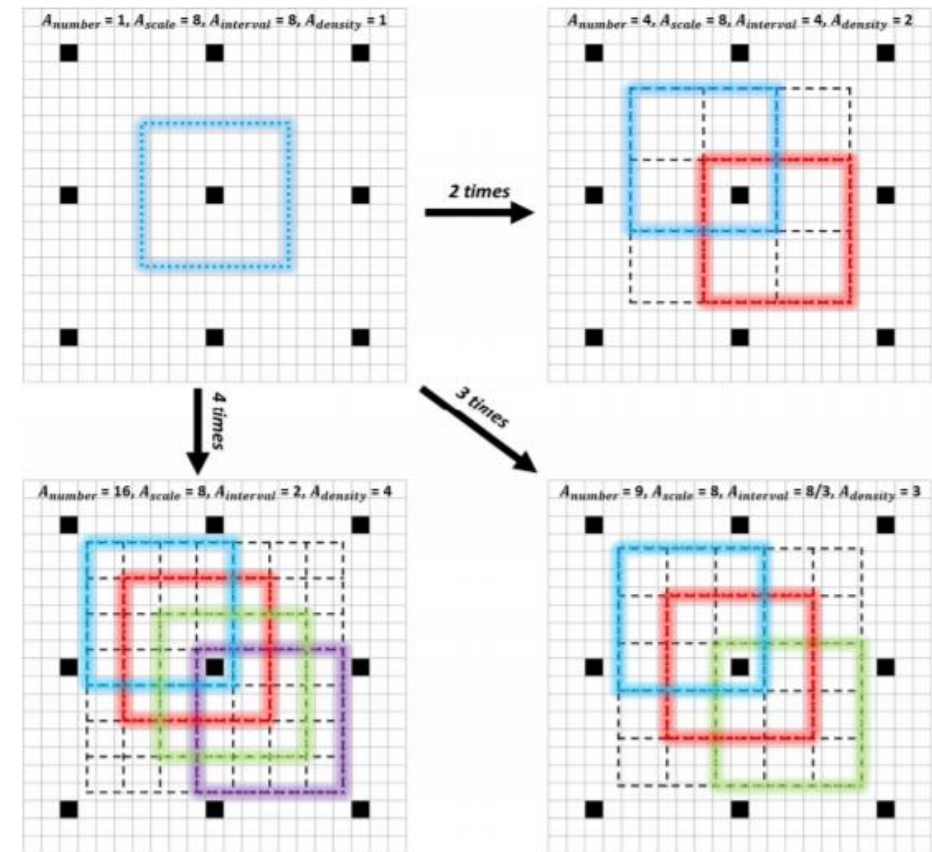
I. Face Detection

B. A Real Example, FaceBoxes [2017-2018]:

Structure: 4. Anchor

- 4.3 min_sizes
[anchor size]:

min_sizes	skips	feature maps
[32x32, 64x64, 128x128]	32	32x32
[256x256]	64	16x16
[512x512]	128	8x8



I. Face Detection

C. Code of FaceBoxes

- **whole structure**
- **prior_box**
- **faceboxes**
- **multibox_loss**

Latest Advanced Code

- Accelerated by CPP
- Structure Refined by SSD
- Added Facial Keypoints

I. Face Detection

D. Advices for Projects

- **BlazeFace**: Google 2019 [Assembled in Google [Mediapipe](#)]
Faster than Face Boxes; real-time in mobile devices
No official code released. **Be No.1 in the world!**
- **FaceBoxes Revised**:
Input size of official FaceBoxes.PyTorch is fixed.
Anchor sizes are fixed.
Let's fix it!

II. Face Recognition

II. Face Recognition

E. What Is Recognition

- **Object**
- **Expression**
- **Face**

II. Face Recognition

E. What Is Recognition

- **Object** **Classification**
- **Expression**
- **Face**

II. Face Recognition

E. What Is Recognition

- **Object** **Classification**
- **Expression** **Classification**
- **Face**

II. Face Recognition

E. What Is Recognition

- **Object** **Classification**
- **Expression** **Classification/Regression**
- **Face**

II. Face Recognition

E. What Is Recognition

- **Object** **Classification** [AffecNet](#)
- **Expression** **Classification/Regression** [Aff-Wild](#)
- **Face**

II. Face Recognition

E. What Is Recognition

- | | | |
|---------------------|---------------------------|---------------------------------|
| • Object | Classification | |
| • Expression | Classification/Regression | <u>AffecNet</u> |
| • Face | “Recognition” | <u>Aff-Wild</u> |

II. Face Recognition

E. What Is Recognition

- **Object** **Classification** [AffecNet](#)
- **Expression** **Classification/Regression** [Aff-Wild](#)
- **Face** **“Recognition”**

Recognition: The procedure of extracting wanted info from data

II. Face Recognition

F. Face Recognition Procedure

- **Object** **Classification** AffecNet
- **Expression** **Classification/Regression** Aff-Wild
- **Face** **“Recognition”**
Let's see this “**real recognition**” in algorithm level

Recognition: The procedure of extracting wanted info from data

II. Face Recognition

E. What Is Recognition

- Object

Classification

AffecNet

- Expression

Classification/Regression

Aff-Wild

- Face

“Recognition”

Let's see this “real recognition” in algorithm level

Recognition
in
Algorithm

Extract features & Retrieve

II. Face Recognition

F. Face Recognition Procedure

F0. Two Steps

Feature Database + Face Recognition

II. Face Recognition

F. Face Recognition Procedure

F1. Construct the Database

0. Face Detection

1. Facial Feature
Extraction

2. Database

II. Face Recognition

F. Face Recognition Procedure

F1. Construct the Database

0. Face Detection

1. Facial Feature
Extraction

2. Database

- Data preparation
- Collected under control

II. Face Recognition

F. Face Recognition Procedure

F1. Construct the Database



- Face represents: Vector
- How to get: VGG / Resnet / ...
- Algorithm: classification

II. Face Recognition

F. Face Recognition Procedure

F1. Construct the Database

0. Face Detection

1. Facial Feature
Extraction

2. Database

➤ Not interested in building

➤ Mad of retrieving

II. Face Recognition

F. Face Recognition Procedure

F2. Face Recognition

0. Face Detection

1. Facial Keypoints
Detection

2. Face Alignment

3. Facial Feature
Extraction

4. Feature Retrieve

II. Face Recognition

F. Face Recognition Procedure

F2. Face Recognition

0. Face Detection

- Most common dataset: 5 / 68 landmarks

1. Facial Keypoints
Detection

- If you're enthusiastic: dlib ([official](#) + [example](#))

2. Face Alignment

3. Facial Feature
Extraction

4. Feature Retrieve



II. Face Recognition

F. Face Recognition Procedure

F2. Face Recognition

0. Face Detection

1. Facial Keypoints
Detection

2. Face Alignment

3. Facial Feature
Extraction

4. Feature Retrieve

➤ 人脸对齐

➤ Remove rotation / light influences / expression / ...

➤ A practical example: [dlib_alignment](#)



II. Face Recognition

F. Face Recognition Procedure

F2. Face Recognition

0. Face Detection




1. Facial Keypoints
Detection

2. Face Alignment

3. Facial Feature
Extraction

4. Feature Retrieve

- Historical: Geometrical Info / Handcrafted Feature
- Now: CNN-based feature (A trained network: VGG / ResNet...)
- A real example: LightCNN
- If you want more: TP-GAN / ... + LightCNN + Perceptual Loss



Feature Vector

II. Face Recognition

F. Face Recognition Procedure

F2. Face Recognition

Pink: Better to know

Green: Good to know

Red: Have to know

Blue: Code

0. Face Detection

➤ Target: Retrieve vector most similar to our feature vector

1. Facial Keypoints
Detection

➤ Problem: How to retrieve fast & accurate

2. Face Alignment

➤ **Methods:** cosine distance (most common way)

3. Facial Feature
Extraction

K-Dimensional Tree (# dimension is small)

PCA (# of dimension is high)

4. Feature Retrieve

cluster / K-nary tree (large scale)

.....

II. Face Recognition


G. A real example: LightCNN (2017)

G1. Framework

- Small — Medium — Large
- We use small as an e.g.

TABLE I
THE ARCHITECTURES OF THE LIGHT CNN-4 MODEL.


Type	Filter Size /Stride	Output Size	#Params
Conv1	$9 \times 9/1$	$120 \times 120 \times 96$	7.7K
MFM1	-	$120 \times 120 \times 48$	-
Pool1	$2 \times 2/2$	$60 \times 60 \times 48$	-
Conv2	$5 \times 5/1$	$56 \times 56 \times 192$	230.4K
MFM2	-	$56 \times 56 \times 96$	-
Pool2	$2 \times 2/2$	$28 \times 28 \times 96$	-
Conv3	$5 \times 5/1$	$24 \times 24 \times 256$	614K
MFM3	-	$24 \times 24 \times 128$	-
Pool3	$2 \times 2/2$	$12 \times 12 \times 128$	-
Conv4	$4 \times 4/1$	$9 \times 9 \times 384$	786K
MFM4	-	$9 \times 9 \times 192$	-
Pool4	$2 \times 2/2$	$5 \times 5 \times 192$	-
fc1	-	512	2,457K
MFM_fc1	-	256	-
Total	-	-	4,095K



II. Face Recognition

G. A real example: LightCNN (2017)

G2. MFM: Max-Feature Map




$$\hat{x}_{ij}^k = \max(x_{ij}^k, x_{ij}^{k+N})$$

II. Face Recognition

G. A real example: LightCNN (2017)

G2. MFM: Max-Feature Map



II. Face Recognition

G. A real example: LightCNN (2017)

G3. Semantic Bootstrapping for Noisy Labels

Firstly, we train the Light CNN-9 model on CASIA-WebFace and fine-tune it on the original noisy labeled MS-Celeb-1M dataset. Second, we employ the trained model to relabel the noisy labeled dataset according to the conditional probabilities $p(t_i|f(x))$. And then we retrain Light CNN-9 on the relabeled dataset. Finally, we further re-sample the original noisy labeled dataset by the retrained model and construct the cleaned MS-Celeb-1M dataset. The details and discussions for

II. Face Recognition

G. A real example: LightCNN (2017)

G3. Semantic Bootstrapping for Noisy Labels

Firstly, we train the Light CNN-9 model on CASIA-WebFace and fine-tune it on the original noisy labeled MS-Celeb-1M dataset. Second, we employ the trained model to relabel the noisy labeled dataset according to the conditional probabilities $p(t_i|f(x))$. And then we retrain Light CNN-9 on the relabeled dataset. Finally, we further re-sample the original noisy labeled dataset by the retrained model and construct the cleaned MS-Celeb-1M dataset. The details and discussions for

- How to deal with bad dataset

II. Face Recognition

G. A real example: LightCNN (2017)

G3. Semantic Bootstrapping for Noisy Labels

Firstly, we train the Light CNN-9 model on CASIA-WebFace and fine-tune it on the original noisy labeled MS-Celeb-1M dataset. Second, we employ the trained model to relabel the noisy labeled dataset according to the conditional probabilities $p(t_i|f(x))$. And then we retrain Light CNN-9 on the relabeled dataset. Finally, we further re-sample the original noisy labeled dataset by the retrained model and construct the cleaned MS-Celeb-1M dataset. The details and discussions for

- **How to deal with bad dataset**

- Train

- Use it

- Train again